

An Automated Cryptanalysis Attack of Two-Time Pads

Implementation and Analysis of the Mason et al. Method

Paul-Constantin Popescu

Faculty of Computer Science, Master in Information Security

Alexandru Ioan Cuza University, Iași, Romania

paul.popescu263@gmail.com

November 4, 2025

Abstract

This paper presents a comprehensive implementation and analysis of the automated cryptanalysis method for two-time pad attacks, as described by Mason et al. in their seminal 2006 ACM CCS paper. When a stream cipher keystream is reused to encrypt multiple plaintexts—a critical security vulnerability—the XOR of the resulting ciphertexts reveals the XOR of the plaintexts themselves. We implement a natural language processing approach using 7-gram character-level language models with Witten-Bell smoothing, combined with the Viterbi algorithm and beam search, to automatically recover both original plaintexts from their XOR. Our implementation achieves 98.7% accuracy on mixed natural language text across 951 test characters, closely matching the original paper’s 99% accuracy on HTML documents. We provide detailed explanations of each algorithmic component, demonstrate the attack’s effectiveness on messages ranging from 49 to 525 bytes, and discuss real-world implications for systems such as Microsoft Word 2002, WEP, and modern AES-CTR implementations.

Keywords: Cryptanalysis, Two-Time Pad, Keystream Reuse, Language Models, Viterbi Algorithm, N-gram Models, Beam Search, Natural Language Processing

1 Introduction

In 1917, Gilbert Vernam introduced the concept of stream ciphers, where a plaintext p is combined with a keystream k via the XOR operation to produce ciphertext: $c = p \oplus k$. Claude Shannon later proved in 1949 that when the keystream is truly random, used only once, and at least as long as the plaintext, this encryption scheme—known as the one-time pad—is theoretically unbreakable [1]. The security proof relies on the fact that without knowledge of the key, every possible plaintext is equally likely, providing perfect secrecy.

However, this theoretical perfection crumbles when a fundamental rule is violated: **keystream reuse**. When the same keystream k is used to encrypt two different plaintexts p and q , the resulting ciphertexts $c_1 = p \oplus k$ and $c_2 = q \oplus k$ can be XORed together to eliminate the keystream:

$$c_1 \oplus c_2 = (p \oplus k) \oplus (q \oplus k) = p \oplus q = x \quad (1)$$

This vulnerability, known as the *two-time pad* problem, has affected numerous real-world cryptographic systems, including Microsoft Word 2002’s RC4 encryption [3], the widely-deployed 802.11 WEP protocol [4], WinZip encryption, and even modern implementations of AES-CTR mode when nonces are improperly reused.

Traditionally, exploiting this vulnerability required manual cryptanalysis: an expert would guess common words or phrases, XOR them with portions of x , and check if the result resembled meaningful text. This process was time-consuming, required significant expertise, and was prone to errors. The breakthrough came in 2006 when Mason et al. [2] published an automated method using natural language processing techniques to systematically recover both plaintexts from their XOR with remarkable accuracy.

1.1 Contributions

This paper makes the following contributions:

1. We provide a complete, faithful implementation of the Mason et al. method in Python, achieving 98.7% accuracy on mixed natural language text—closely matching the original paper’s 99% accuracy on HTML documents.
2. We present a comprehensive, line-by-line explanation of the algorithm, mapping each code component to the corresponding sections of the original paper, making the method accessible to students and practitioners.
3. We demonstrate the attack’s effectiveness across multiple message lengths (49 to 525 bytes) and analyze where errors occur, providing insights into the method’s strengths and limitations.
4. We discuss the real-world implications of keystream reuse and provide guidance on how modern cryptographic implementations can avoid this vulnerability.

The remainder of this paper is organized as follows: Section 2 provides background on stream ciphers and the two-time pad problem; Section 3 details the cryptanalysis methodology; Section 4 describes our implementation; Section 5 presents experimental results; Section 6 discusses implications and limitations; and Section 7 concludes.

2 Background and Related Work

2.1 Stream Ciphers and the One-Time Pad

A stream cipher generates a pseudorandom keystream from a secret key and combines it with the plaintext using the XOR operation. The security of such systems depends critically on the properties of the keystream. Shannon’s perfect secrecy theorem states that the one-time pad provides unconditional security when three conditions are met: (1) the key is truly random, (2) the key is at least as long as the message, and (3) **the key is never reused**.

In practice, true randomness and key distribution are challenging, leading to the widespread use of pseudorandom stream ciphers like RC4, ChaCha20, and AES in counter (CTR) mode. These ciphers generate keystreams from shorter keys using cryptographic algorithms, sacrificing perfect secrecy for practical usability.

2.2 The Two-Time Pad Vulnerability

The two-time pad problem arises when the same keystream encrypts multiple messages. As shown in Equation 1, the keystream cancels out when ciphertexts are XORed, leaving only the XOR of the plaintexts. This significantly reduces the attacker’s search space: instead of searching through all possible keys, the attacker need only find plaintext pairs (p, q) that satisfy $p \oplus q = x$ and resemble natural language.

For each position i in the XOR stream, approximately 95 character pairs satisfy the constraint $p_i \oplus q_i = x_i$ (considering printable ASCII characters). For a message of length ℓ , this yields approximately 95^ℓ possible solutions. While still astronomically large for even modest message lengths (e.g., $95^{100} \approx 10^{199}$), this is vastly smaller than the 256^ℓ possibilities in the general case.

2.3 Prior Cryptanalysis Approaches

Early attempts at automated two-time pad cryptanalysis were limited. Dawson and Nielsen [5] achieved 62–76% accuracy using dictionary-based methods and simulated annealing, but their approach required significant computational resources and manual intervention.

Historical examples of manual two-time pad cryptanalysis include the VENONA project, where NSA cryptanalysts spent decades decrypting Soviet diplomatic communications that had reused one-time pad keys due to wartime key shortage. This laborious process required expert cryptanalysts and yielded only partial decryptions.

2.4 The Mason et al. Approach

Mason et al. [2] revolutionized two-time pad cryptanalysis by applying natural language processing techniques. Their key insights were:

1. **Statistical language models** can distinguish real text from random character sequences by modeling character transition probabilities.
2. **Dynamic programming** (specifically, the Viterbi algorithm) can efficiently find the most probable plaintext pair without exhaustive search.
3. **Beam search** makes the algorithm tractable by pruning low-probability paths, reducing the exponential state space to manageable size.

Their method achieved 99% accuracy on HTML documents and demonstrated practical attacks against real systems, including Microsoft Word 2002’s encryption scheme.

3 Methodology

The cryptanalysis method consists of four main components: (1) language model construction, (2) probability smoothing, (3) Viterbi decoding with beam search, and (4) stream switching detection. We describe each in detail.

3.1 N-gram Language Models

A character-level n -gram model predicts the probability of a character based on the $n - 1$ preceding characters. Formally, for a text sequence $c_1 c_2 \dots c_\ell$, the model estimates:

$$P(c_i | c_{i-n+1}, c_{i-n+2}, \dots, c_{i-1}) \quad (2)$$

Following Mason et al., we use $n = 7$ (7-grams), which considers the previous six characters when predicting the next one. This provides sufficient context to capture common word patterns and phrases while maintaining tractable training data requirements.

3.1.1 Training Process

The training process constructs three essential data structures from a text corpus:

1. `ngram_counts`: Maps each (context, character) pair to its frequency count. For example, `ngram_counts[("hello", " ")]` stores how many times a space follows "hello" in the training corpus.
2. `context_counts`: Maps each context to its total occurrence count. This serves as the denominator for probability calculations.
3. `unique_continuations`: Maps each context to the set of unique characters that have followed it. This is used in Witten-Bell smoothing (Section 3.2).

We also incorporate special Begin-Of-Message (BOM) and End-Of-Message (EOM) markers, represented as `\x00` and `\x01` respectively. These markers help the model predict the first and last characters more accurately, as messages often begin with capital letters and end with punctuation.

The training algorithm processes the corpus character by character, counting all n -grams from 1-gram up to 7-gram. This multi-scale counting is essential for the backoff mechanism described in the next section.

3.2 Witten-Bell Smoothing

A critical challenge in language modeling is handling unseen n -grams—character sequences that never appeared in the training corpus. Assigning zero probability to such sequences would cause the Viterbi algorithm to fail when encountering them.

Witten-Bell smoothing [6] addresses this through a principled interpolation between different context lengths. The probability of character c following context h is computed as:

$$P(c|h) = \begin{cases} (1 - \lambda(h)) \cdot \frac{N(hc)}{N(h)} + \lambda(h) \cdot P(c|h') & \text{if } N(hc) > 0 \\ \lambda(h) \cdot P(c|h') & \text{otherwise} \end{cases} \quad (3)$$

where:

- $N(hc)$ is the count of context h followed by character c
- $N(h)$ is the total count of context h
- h' is context h with the first character removed (backoff context)
- $\lambda(h) = \frac{T(h)}{T(h)+N(h)}$ is the smoothing parameter
- $T(h)$ is the number of unique characters that have followed h

The smoothing parameter $\lambda(h)$ balances trust between the current context and shorter contexts. When a context has been seen many times with few unique continuations, $\lambda(h)$ is small, and we trust the direct probability estimate. When a context is rare or has many unique continuations, $\lambda(h)$ is larger, and we rely more on the backoff model.

The recursion terminates at the base case of an empty context, where we assume a uniform distribution over the 95 printable ASCII characters:

$$P(c|\emptyset) = \frac{1}{95} \quad (4)$$

3.3 The Viterbi Algorithm with Beam Search

The Viterbi algorithm is a dynamic programming method for finding the most probable sequence in a probabilistic model. Originally developed for decoding convolutional codes and widely used in speech recognition, we adapt it here for two-time pad cryptanalysis.

3.3.1 State Space

A state at position i in the message is represented as a triple:

$$s_i = (i, \text{ctx}_1, \text{ctx}_2) \quad (5)$$

where ctx_1 and ctx_2 are the last $n - 1$ characters of plaintexts p and q respectively. For 7-grams, each context contains six characters.

The initial state at position 0 is:

$$s_0 = (0, \text{BOM}, \text{BOM}) \quad (6)$$

with probability 1.0 (or log-probability 0.0).

3.3.2 Transitions

From each state $s_i = (i, \text{ctx}_1, \text{ctx}_2)$ with accumulated log-probability L_i , we consider all character pairs (p_i, q_i) that satisfy:

$$p_i \oplus q_i = x_i \quad (7)$$

where x_i is the i -th byte of the XOR stream, and both p_i and q_i are printable ASCII characters (32–126).

For each valid pair, we compute the transition probability using our language models:

$$P(p_i, q_i | \text{ctx}_1, \text{ctx}_2) = P(p_i | \text{ctx}_1) \cdot P(q_i | \text{ctx}_2) \quad (8)$$

This assumes independence between the two plaintext streams, which is reasonable for most practical scenarios.

The new state becomes:

$$s_{i+1} = (i + 1, \text{ctx}'_1, \text{ctx}'_2) \quad (9)$$

where ctx'_1 and ctx'_2 are updated by appending the new characters and keeping only the last $n - 1$ characters.

The accumulated log-probability is:

$$L_{i+1} = L_i + \log P(p_i | \text{ctx}_1) + \log P(q_i | \text{ctx}_2) \quad (10)$$

3.3.3 Viterbi Optimization

Multiple paths may lead to the same state. The Viterbi algorithm keeps only the path with the highest probability, discarding all others. This is optimal because:

Theorem 3.1 (Optimal Substructure). *If the optimal path to position ℓ passes through state s_i at position i , then the segment from position 0 to i must be the optimal path to state s_i .*

This property allows us to prune suboptimal paths early, preventing exponential explosion of the search space.

3.3.4 Beam Search

Even with Viterbi optimization, the number of states can grow exponentially with message length. At position i , there could be up to $95^{2(n-1)} \approx 10^{23}$ possible states for 7-grams (each context can be any string of 6 printable characters).

Beam search addresses this by keeping only the B most probable states at each position, where B is the beam width. States are sorted by log-probability, and only the top B are retained for further expansion.

The beam width represents a trade-off:

- **Larger beam:** Higher accuracy, slower computation, more memory
- **Smaller beam:** Faster computation, less memory, risk of pruning the optimal solution

Mason et al. recommend beam widths of 100–500, depending on message length and accuracy requirements. Our experiments use beam widths ranging from 100 (for 49-byte messages) to 400 (for 525-byte messages).

3.3.5 Algorithm Summary

Algorithm 1 presents the complete Viterbi decoding procedure with beam search.

Algorithm 1 Viterbi Decoding with Beam Search

1: **Input:** XOR stream x of length ℓ , language models M_1 and M_2 , beam width B
2: **Output:** Recovered plaintexts p and q

3:
4: $\text{states} \leftarrow \{(0, \text{BOM}, \text{BOM}) : (0.0, [])\}$
5: **for** $i = 0$ to $\ell - 1$ **do**
6: $\text{new_states} \leftarrow \emptyset$
7: **for each** $(pos, \text{ctx}_1, \text{ctx}_2), (L, \text{path})$ in states **do**
8: **if** $pos \neq i$ **then**
9: **continue**
10: **end if**
11: **for** $p_i = 32$ to 126 **do**
12: $q_i \leftarrow p_i \oplus x_i$
13: **if** $q_i < 32$ or $q_i > 126$ **then**
14: **continue**
15: **end if**
16: $P_1 \leftarrow M_1.\text{probability}(p_i, \text{ctx}_1)$
17: $P_2 \leftarrow M_2.\text{probability}(q_i, \text{ctx}_2)$
18: **if** $P_1 < 10^{-20}$ or $P_2 < 10^{-20}$ **then**
19: **continue**
20: **end if**
21: $L_{\text{new}} \leftarrow L + \log(P_1) + \log(P_2)$
22: $\text{ctx}'_1 \leftarrow (\text{ctx}_1 + p_i)[-6 :]$
23: $\text{ctx}'_2 \leftarrow (\text{ctx}_2 + q_i)[-6 :]$
24: $s_{\text{new}} \leftarrow (i + 1, \text{ctx}'_1, \text{ctx}'_2)$
25: $\text{path}_{\text{new}} \leftarrow \text{path} + [(p_i, q_i)]$
26: **if** $s_{\text{new}} \notin \text{new_states}$ or $L_{\text{new}} > \text{new_states}[s_{\text{new}}][0]$ **then**
27: $\text{new_states}[s_{\text{new}}] \leftarrow (L_{\text{new}}, \text{path}_{\text{new}})$
28: **end if**
29: **end for**
30: **end for**
31: $\text{states} \leftarrow \text{top_}_B(\text{new_states})$ ▷ Keep only top B by log-probability
32: **end for**
33:
34: $(L_{\text{best}}, \text{path}_{\text{best}}) \leftarrow \max_{s \in \text{states}} \text{states}[s][0]$
35: $p \leftarrow \text{path}_{\text{best}}[:, 0]$ ▷ Extract first element of each pair
36: $q \leftarrow \text{path}_{\text{best}}[:, 1]$ ▷ Extract second element of each pair
37: **return** p, q

3.4 Stream Switching Detection

A subtle problem can occur during decoding: the algorithm may correctly identify character pairs but assign them to the wrong plaintexts. This happens because after an error, the algorithm can "lose track" of which context corresponds to which plaintext stream.

For example, the algorithm might correctly decode pairs $(H, S), (e, e), (l, c), (l, r), (o, e), (t, t)$ but incorrectly produce:

- $p = \text{"Secret"}$
- $q = \text{"Hello"}$

instead of the correct:

- $p = \text{"Hello"}$
- $q = \text{"Secret"}$

The pairs are correct (verified by checking $p \oplus q = x$), but the assignment is wrong.

Mason et al. propose a post-processing solution: after decoding, score both possible assignments using the language models:

$$S_{\text{normal}} = \log P_{M_1}(p) + \log P_{M_2}(q) \quad (11)$$

$$S_{\text{swapped}} = \log P_{M_1}(q) + \log P_{M_2}(p) \quad (12)$$

If $S_{\text{swapped}} > S_{\text{normal}}$, the streams were likely swapped, and we exchange p and q .

This technique is particularly effective when the two language models are trained on distinctly different text types (e.g., formal vs. casual), as the models can distinguish which plaintext belongs to which genre.

4 Implementation

We implemented the complete cryptanalysis system in Python 3, totaling approximately 300 lines of code. The implementation consists of two main classes: `LanguageModel` and `TwoTimePadCracker`.

4.1 Language Model Implementation

The `LanguageModel` class encapsulates all functionality related to n-gram probability computation. Key methods include:

- `__init__(n=7)`: Initializes data structures for n-gram counting
- `train(corpus)`: Processes a text corpus and builds the probability model
- `get_probability(char, context)`: Returns $P(\text{char}|\text{context})$ using Witten-Bell smoothing
- `score_text(text)`: Computes total log-probability of a text sequence

The implementation uses Python's `defaultdict` for efficient storage of sparse n-gram counts. Listing 1 shows the core probability calculation with smoothing.

```
1 def _probability_recursive(self, char, context):  
2     # Base case: empty context  
3     if len(context) == 0:  
4         return 1.0 / 95    # Uniform distribution  
5
```

```

6     context_total = self.context_counts.get(context, 0)
7
8     # Backoff if context never seen
9     if context_total == 0:
10        return self._probability_recursive(char, context[1:])
11
12     # Witten-Bell smoothing
13     unique_chars = len(self.unique_continuations[context])
14     count = self.ngram_counts.get((context, char), 0)
15     lambda_param = unique_chars / (context_total + unique_chars)
16
17     if count > 0:
18        prob_direct = count / context_total
19        prob_backoff = self._probability_recursive(char, context[1:])
20        return (1 - lambda_param) * prob_direct + \
21               lambda_param * prob_backoff
22     else:
23        return lambda_param * self._probability_recursive(char, context[1:])
24

```

Listing 1: Witten-Bell Smoothing Implementation

4.2 Two-Time Pad Cracker Implementation

The `TwoTimePadCracker` class implements the Viterbi algorithm with beam search. The main method is `crack(xor_stream)`, which returns the recovered plaintext pair.

Key implementation details include:

1. **State Representation:** States are tuples `(position, context1, context2)`, stored in a dictionary mapping to `(log_probability, path)`.
2. **Logarithmic Computation:** All probabilities are converted to log-space immediately to prevent underflow. We use Python’s `math.log()` function.
3. **Beam Pruning:** After expanding all states at position i , we sort by log-probability and keep only the top B states using Python’s built-in `sorted()` function.
4. **Progress Reporting:** The implementation includes verbose mode that prints progress every 100 positions, including estimated time to completion.

4.3 Training Corpus Generation

Following Mason et al.’s guidance, we constructed two training corpora representing different text types:

1. **Business/Formal Corpus** (1.37M characters): Professional emails, meeting notifications, business correspondence, and formal documentation
2. **Secret/Casual Corpus** (1.62M characters): Casual messages, secret communications, personal notes, and informal text

Each corpus was generated using data augmentation techniques:

- 56–64 base sentences representing typical phrases
- 40 repetitions with variations (lowercase, spacing, word reorderings)

- Addition of common English phrases for baseline coverage

This approach generated sufficient training data while maintaining control over the text types and distributions. The resulting corpora contain 15,295 and 19,757 unique 7-gram contexts respectively, providing good coverage of common character sequences.

4.4 Computational Complexity

The implementation’s time complexity is:

$$O(\ell \cdot B \cdot A \cdot \log B) \quad (13)$$

where ℓ is message length, B is beam width, $A \approx 95$ is the alphabet size, and the $\log B$ factor comes from sorting states during beam pruning.

Space complexity is dominated by the language model storage:

$$O(|C| \cdot n) \quad (14)$$

where $|C|$ is corpus size and $n = 7$ is the n-gram size. For our 3M character corpus and our hardware configuration (Intel i7-12700KF with 16 GB RAM), this requires approximately 200 MB of RAM, which is easily manageable.

4.5 Software Dependencies

Unlike the original Mason et al. implementation, which used the LingPipe Java toolkit [7] for language modeling functionality, our implementation requires no external NLP libraries. We implemented all language modeling components—n-gram counting, Witten-Bell smoothing, and probability calculation—from first principles using only Python’s standard library.

This design choice offers several advantages:

- **Transparency:** Every algorithmic detail is explicit in the code, making it easier to understand and verify correctness.
- **Portability:** No external dependencies means the implementation runs anywhere Python is available.
- **Educational Value:** Students can see exactly how n-gram models and smoothing work without black-box library calls.
- **Customizability:** Easy to modify smoothing parameters, add new features, or experiment with variations.

The only Python packages used are from the standard library: `collections.defaultdict` for efficient sparse data structures, `math` for logarithmic calculations, and `time` for performance measurement.

5 Experimental Results

We evaluated our implementation on four test cases of increasing length, designed to represent realistic scenarios where keystream reuse might occur.

5.1 Test Configuration

All experiments were conducted on a system with:

- CPU: Intel Core i7-12700KF (12 cores, 3.6 GHz base clock)
- RAM: 16 GB DDR5

- Motherboard: Gigabyte B760M GAMING X AX
- OS: Microsoft Windows 11 Pro (Build 22631)
- Python: 3.12

Beam widths were selected based on message length:

- 49 bytes: $B = 100$
- 126 bytes: $B = 200$
- 251 bytes: $B = 300$
- 525 bytes: $B = 400$

5.2 Quantitative Results

Table 1 summarizes the accuracy and performance across all test cases.

Table 1: Experimental Results on Four Test Cases

Test Case	Length (bytes)	Beam Width	Accuracy (%)	Time (seconds)
Short Message	49	100	87.8	1.66
Medium Message	126	200	96.0	10.05
Long Message	251	300	100.0	33.93
Very Long Message	525	400	99.8	112.49
Overall	951	—	98.7	158.12

The overall accuracy of 98.7% represents 939 correctly recovered characters out of 951 total across all test cases. Notably, the 251-byte test achieved perfect 100% accuracy, and the 525-byte test had only a single error at the final position.

5.3 Performance Characteristics

Processing speed averaged 166.3 ms per byte, which compares favorably to Mason et al.’s reported 200 ms/byte on 2006-era hardware. The performance improvement is primarily attributable to advances in processor technology over the 18-year span.

Figure ?? illustrates the relationship between message length and processing time, showing approximately linear scaling as predicted by the complexity analysis.

5.4 Error Analysis

Of the 12 total errors across all test cases:

- **8 errors** (67%) occurred within the last 5 characters of messages
- **3 errors** (25%) involved rare character combinations not well-represented in training
- **1 error** (8%) occurred mid-message, likely due to beam pruning

The concentration of errors at message boundaries is expected, as these positions have limited context for prediction. The BOM and EOM markers help, but cannot fully compensate for the reduced context window.

Table 2: Comparison with Original Paper Results

Metric	Mason et al. (2006)	Our Implementation
Accuracy on HTML	99%	—
Accuracy on English	95–97%	98.7%
Processing Speed	200 ms/byte	166 ms/byte
Training Corpus Size	600K chars	3M chars
Unique Contexts	Not reported	15K–20K
Beam Width	100–500	100–400

5.5 Comparison to Original Paper

Table 2 compares our results to those reported by Mason et al.

Our 98.7% accuracy on mixed natural language text is comparable to the paper’s 99% on HTML documents. This is particularly noteworthy because HTML has more predictable structure (tags, attributes) than free-form natural language, making it an easier target for language models.

5.6 Stream Switching Detection Effectiveness

The stream switching detection mechanism correctly identified the proper plaintext assignment in all four test cases. The score differences were substantial:

- Test 1: $S_{\text{normal}} = -230.06$ vs. $S_{\text{swapped}} = -608.12$ (difference: 378)
- Test 2: $S_{\text{normal}} = -549.55$ vs. $S_{\text{swapped}} = -1454.62$ (difference: 905)
- Test 3: $S_{\text{normal}} = -741.72$ vs. $S_{\text{swapped}} = -3056.29$ (difference: 2314)
- Test 4: $S_{\text{normal}} = -1109.05$ vs. $S_{\text{swapped}} = -6504.72$ (difference: 5396)

The large score differences indicate that the language models strongly distinguish between formal and casual text styles, making stream assignment unambiguous.

6 Discussion

6.1 Real-World Implications

The high accuracy achieved by our implementation demonstrates that two-time pad attacks are not merely theoretical—they are practical, automated, and highly effective. This has serious implications for systems that reuse keystreams.

6.1.1 Historical and Current Vulnerabilities

Several real-world systems have been vulnerable to two-time pad attacks:

1. **Microsoft Word 2002:** Wu [3] demonstrated that Word’s RC4 encryption reused initialization vectors when documents were edited and re-saved, enabling two-time pad attacks on document revisions.
2. **802.11 WEP:** The Wired Equivalent Privacy protocol used 24-bit initialization vectors with RC4, guaranteeing collisions (and thus keystream reuse) after approximately $2^{12} \approx 4000$ packets due to the birthday paradox [4].
3. **WinZip Encryption:** Older versions used a proprietary stream cipher with inadequate key derivation, allowing keystream reuse across files encrypted with the same password.

4. **AES-CTR Mode:** While AES-CTR is secure when used correctly, implementations that fail to ensure nonce uniqueness are vulnerable to exactly this attack.

6.1.2 Modern Relevance

Despite being published in 2006, this attack remains relevant today:

- **Legacy Systems:** Many organizations still use older software with known keystream reuse vulnerabilities.
- **Implementation Errors:** Even with modern algorithms like AES-GCM and ChaCha20-Poly1305, implementation bugs can lead to nonce reuse.
- **IoT Devices:** Resource-constrained devices sometimes take shortcuts that compromise security, including improper key management.

6.2 Defense Mechanisms

Preventing two-time pad attacks requires adherence to fundamental cryptographic principles:

1. **Never Reuse Keystreams:** This is the cardinal rule. Every encryption operation must use a unique keystream.
2. **Use Proper Nonces/IVs:**
 - Generate cryptographically random IVs
 - Use counters that never repeat
 - Ensure sufficient IV/nonce space (128 bits recommended)
3. **Employ Authenticated Encryption:** Modern AEAD schemes (AES-GCM, ChaCha20-Poly1305, AES-GCM-SIV) detect tampering and often include safeguards against nonce misuse.
4. **Key Rotation:** Regularly rotate encryption keys to limit exposure from any single key compromise.
5. **Use High-Level Cryptographic Libraries:** Libraries like libsodium and NaCl provide misuse-resistant interfaces that make it difficult to accidentally reuse keystreams.

6.3 Limitations of the Method

While highly effective on natural language text, the method has limitations:

1. **Requires Text Structure:** The attack fails on compressed, encrypted, or truly random data, as these lack the statistical patterns that language models exploit.
2. **Needs Representative Training Data:** The language models must be trained on text similar to the target plaintexts. Highly specialized jargon or non-standard language may reduce accuracy.
3. **Computational Cost:** While tractable, processing very long messages (thousands of bytes) with large beam widths becomes expensive. Our 525-byte test took nearly two minutes.
4. **Non-Optimal with Small Beams:** Beam search is a heuristic. With insufficient beam width, the optimal solution may be pruned, reducing accuracy.
5. **Assumes Independence:** The method assumes the two plaintexts are independent. If they are related (e.g., two versions of the same document with minor edits), the independence assumption breaks down, though the attack may still succeed.

6.4 Potential Improvements

Several enhancements could further improve the method:

1. **Neural Language Models:** Modern transformer-based models (e.g., GPT-style architectures) could provide better probability estimates than n-grams, potentially improving accuracy.
2. **Adaptive Beam Width:** Dynamically adjust beam width based on local ambiguity—use larger beams in uncertain regions, smaller beams in confident regions.
3. **Iterative Refinement:** After initial decoding, use the recovered text to retrain or adapt the language models, then re-decode with improved models.
4. **Multiple Keystream Reuses:** If the same keystream encrypted three or more plaintexts, additional constraints could significantly improve accuracy.
5. **Parallel Processing:** Beam search is embarrassingly parallel—different beams can be processed independently, enabling GPU acceleration.

6.5 Ethical Considerations

This research has potential for both beneficial and malicious use:

Beneficial Applications:

- Security auditing and penetration testing
- Historical cryptanalysis (e.g., analyzing VENONA transcripts)
- Educational demonstrations of cryptographic vulnerabilities
- Motivating proper cryptographic implementation

Potential Misuse:

- Unauthorized decryption of communications
- Privacy violations
- Industrial espionage

We emphasize that this work is intended for educational and defensive purposes only. Unauthorized access to encrypted communications is illegal in most jurisdictions and ethically indefensible.

7 Conclusion

We have presented a complete implementation and comprehensive analysis of the Mason et al. method for automated two-time pad cryptanalysis. Our implementation achieves 98.7% accuracy on mixed natural language text, closely matching the original paper’s results and demonstrating that the attack remains highly effective nearly two decades after its publication.

The success of this method underscores a fundamental lesson in cryptography: *theoretical security guarantees evaporate when basic operational rules are violated*. The one-time pad is provably unbreakable—but only when the keystream is never reused. In practice, keystream reuse has affected numerous real-world systems, from Soviet diplomatic communications to modern wireless protocols.

Several key insights emerge from this work:

1. **Natural language structure enables cryptanalysis:** The statistical patterns in human language, captured by n-gram models, provide sufficient information to distinguish correct decryptions from random character sequences.

2. **Modern NLP techniques transfer to cryptanalysis:** Algorithms developed for speech recognition and natural language processing—specifically, the Viterbi algorithm with beam search—prove highly effective for breaking cryptographic protocols.
3. **The attack is practical:** With 98.7% accuracy and processing speeds of 166 ms/byte, the attack is fast enough for real-world application against systems with keystream reuse vulnerabilities.
4. **Prevention is straightforward but critical:** The attack is completely prevented by ensuring keystream uniqueness, emphasizing the importance of proper cryptographic implementation.

For practitioners, the message is clear: cryptographic protocols must be implemented correctly, with particular attention to key and nonce management. For researchers, this work demonstrates the power of interdisciplinary approaches, combining cryptography, natural language processing, and dynamic programming to solve challenging problems.

Future work could explore extensions to other scenarios (e.g., partial keystream reuse, corrupted ciphertexts), application of modern neural language models, and development of more sophisticated training corpus generation techniques. Additionally, investigating the method’s effectiveness on non-English languages and specialized technical text would broaden its applicability.

The code and implementation details are available at <https://github.com/yourusername/two-time-pad-crack> for educational purposes, encouraging further experimentation and improvement by the research community.

Acknowledgements

The author thanks Professor [Name] for guidance on this project, colleagues [Names] for valuable discussions, and the authors of the original Mason et al. paper for their groundbreaking work. This research was conducted as part of the Master’s program in Information Security at Alexandru Ioan Cuza University.

References

- [1] Shannon, C.E. (1949). Communication theory of secrecy systems. *Bell System Technical Journal*, 28(4), 656–715.
- [2] Mason, J., Watkins, K., Eisner, J., & Stubblefield, A. (2006). A natural language approach to automated cryptanalysis of two-time pads. In *Proceedings of the 13th ACM Conference on Computer and Communications Security* (CCS ’06), 235–244.
- [3] Wu, H. (2005). The misuse of RC4 in Microsoft Word and Excel. *Cryptology ePrint Archive*, Report 2005/007.
- [4] Borisov, N., Goldberg, I., & Wagner, D. (2001). Intercepting mobile communications: The insecurity of 802.11. In *Proceedings of the 7th Annual International Conference on Mobile Computing and Networking* (MobiCom ’01), 180–189.
- [5] Dawson, E., & Nielsen, L. (1996). Automated cryptanalysis of XOR plaintext strings. *Cryptologia*, 20(2), 165–181.
- [6] Witten, I.H., & Bell, T.C. (1991). The zero-frequency problem: Estimating the probabilities of novel events in adaptive text compression. *IEEE Transactions on Information Theory*, 37(4), 1085–1094.
- [7] Alias-i. (2008). LingPipe 4.1.0. Available from <http://alias-i.com/lingpipe>